



MCMC for Cut Models or Chasing a Moving Target with MCMC

Martyn Plummer

International Agency for Research on Cancer

MCMSki

Chamonix, 6 Jan 2014

Closed-form solution

This simple example has a closed form solution.
The density of θ is given by the mixture:

$$p^*(\theta) = \int p(\theta | z) \phi(z) dz$$


Taking the limit of an improper, flat prior on θ

$$\theta | z \sim N(z, 1)$$

$$z \sim N(0, 1)$$

$$\theta \sim N(0, 2)$$

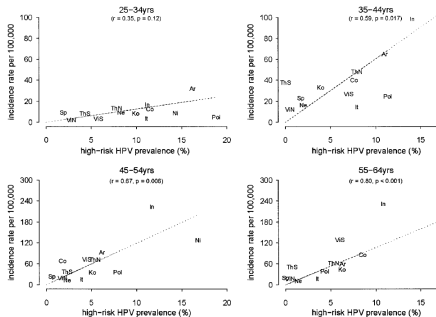
Results of naive cut algorithm by sampling method ¹



Different sampling algorithms converge to different limiting distributions
Only the conjugate sampler is correct.

Cuts in measurement error models: Motivating example

There is an ecological association between HPV prevalence and cervical cancer incidence².



HPV is a necessary cause of cancer, but risk is modulated by other cofactors: smoking, childbirth, hormonal contraceptives, ...

International Agency for Research on Cancer




A toy measurement error model for the ecological data

We experimented with a functional measurement error model for these data, with a Poisson regression model for incidence and a binomial model for (age-specific) prevalence:

$$\begin{array}{lll} Y_i & \sim & \text{Poisson}(N_i \exp(\lambda_i)) \quad \text{Cancer incidence data} \\ \lambda_i & = & \theta_1 + \theta_2 \varphi_i \quad \text{Incidence rates} \\ Z_i & \sim & \text{Bin}(n_i, \varphi_i) \quad \text{HPV prevalence data} \end{array}$$

A cut measurement error model



In a cut model, the graph G is divided into two sub-graphs G_1, G_2 .

- ▶ Nodes in G_1 are updated ignoring nodes in G_2 .
- ▶ Nodes in G_2 are updated as normal (naive cut algorithm).


- ▶ In our example, we use only prevalence survey data (Z) to estimate HPV prevalence (φ)
- ▶ Feedback from cancer incidence rates (Y) via the putative dose-response relationship is cut.

Definition of cut model in BUGS

OpenBUGS provides a cut function to denote cuts, and implements the naive cut algorithm when sampling cut models.

```
model {  
  
  ## Disease model  
  for (j in 1:13) {  
    ncases[j] ~ dpois(mean[j])  
    log(mean[j]) <- theta[1] + phi.cut[j] * theta[2] + log(Npop[j] * 1.0E-3)  
  }  
  
  theta[1] ~ dnorm(0, 1.0E-3)  
  theta[2] ~ dnorm(0, 1.0E-3)  
  
  ## Cuts  
  for (j in 1:13) {  
    p.cut[j] <- cut(p[j])  
  }  
  
  ## Measurement model - below the cut  
  for (j in 1:13) {  
    npositives[j] ~ dbin(phi[j], Nsubjects[j])  
  }  
  
  ## Exposure model - below the cut  
  for (j in 1:13) {  
    phi[j] ~ dunif(0, 1)  
  }  
}
```

Results of naive cut algorithm for θ_2 by sampling method ³



Different update methods converge to different limiting distributions. The correct distribution could be calculated by multiple imputation, but is not shown here.

Why the naive cut algorithm does not work (2/2)

In general, MCMC methods do not sample directly from the target density but supply a reversible transition $\theta^{t-1} \rightarrow \theta^t$ at iteration t . The transition is in detailed balance with the full conditional distribution:

$$p(\theta^{t-1} \mid \mathbf{Y}, \varphi^t) p(\theta^{t-1} \rightarrow \theta^t \mid \varphi^t) = \\ p(\theta^t \mid \mathbf{Y}, \varphi^t) p(\theta^t \rightarrow \theta^{t-1} \mid \varphi^t)$$

But for $p^*(\theta)$ to be the stationary distribution we need:

$$p(\theta^{t-1} \mid \mathbf{Y}, \varphi^{t-1}) p(\theta^{t-1} \rightarrow \theta^t \mid \varphi^{t-1}, \varphi^t) = \\ p(\theta^t \mid \mathbf{Y}, \varphi^t) p(\theta^t \rightarrow \theta^{t-1} \mid \varphi^t, \varphi^{t-1})$$

Why the naive cut algorithm does not work (2/2)

In general, MCMC methods do not sample directly from the target density but supply a reversible transition $\theta^{t-1} \rightarrow \theta^t$ at iteration t . The transition is in detailed balance with the full conditional distribution:

$$p(\theta^{t-1} \mid \mathbf{Y}, \varphi^t) p(\theta^{t-1} \rightarrow \theta^t \mid \varphi^t) = p(\theta^t \mid \mathbf{Y}, \varphi^t) p(\theta^t \rightarrow \theta^{t-1} \mid \varphi^t)$$

But for $p^*(\theta)$ to be the stationary distribution we need:

$$p(\theta^{t-1} \mid \mathbf{Y}, \varphi^{t-1}) p(\theta^{t-1} \rightarrow \theta^t \mid \varphi^{t-1}, \varphi^t) = p(\theta^t \mid \mathbf{Y}, \varphi^t) p(\theta^t \rightarrow \theta^{t-1} \mid \varphi^t, \varphi^{t-1})$$

The balance relation uses the **current and previous** values of φ .

Can we modify a standard MCMC update? (1/2)

Maybe we can add a Metropolis-Hastings acceptance step, treating the move $\theta^{t-1} \rightarrow \theta^t$ as a proposal to be accepted with probability $\min(1, R)$ where

$$R = \frac{p(\theta^t \mid \mathbf{Y}, \varphi^t) p(\theta^t \rightarrow \theta^{t-1} \mid \varphi^{t-1})}{p(\theta^{t-1} \mid \mathbf{Y}, \varphi^{t-1}) p(\theta^{t-1} \rightarrow \theta^t \mid \varphi^t)}$$

Note that $R = 1$ in the case of direct sampling:

$$p(\theta^{t-1} \rightarrow \theta^t \mid \varphi) = p(\theta^t \mid \mathbf{Y}, \varphi)$$

Reweighting the samples (latest attempt cont.)

- ▶ This MCMC process generates a sequence of samples with the correct conditional distribution $\theta \mid \varphi$ but the wrong marginal distribution for φ .
- ▶ It must be reweighted:
 - ▶ Give weight $1/n_t$ to each of $\theta^{t1} \dots \theta^{tn_t}$

